SEO Multi-Agent Tool

Non-Tech Guide

For business, marketing, and operations. How to generate content on the web page and via Google Sheet.

Who this is for: Business, marketing, and operations people
Language: Simple English only

Engineers should also open Technical.md / Technical.pdf.

1. Why we built this (the real purpose)

We are not here for likes and views

Many websites and social posts are written mainly to get:

- more views
- more followers
- more attention

Those posts can look exciting. People click. People scroll. But after reading, the person still feels stuck.

That is not our goal.

We are here to help people with real problems

We create content so readers can learn something useful and feel less stuck with a real issue - in any topic or domain.

If our content cannot fully solve their problem, it should still:

- point them clearly to where to go next, and
- do that without confusing them

Simple rule: help the reader -> not chase attention.

Example (easy to understand)

A person wants to learn LangGraph and persistent memory.

They search Google or social media. They find many popular posts with lots of views. Those posts may:

- sound exciting for a minute
- use big words
- skip the hard but important parts

So the person feels interested... but still does not really understand, and still cannot fix their issue.

Our articles and blogs should fill that gap.

After reading our content, the person should either:

1. Actually learn something they can use, or
2. Know the clear next step to learn it properly - without feeling lost

That is why this whole workflow exists.

2. What is this tool?

It is a content helper.

Every morning (example: 9 AM), an automation tool called n8n looks at a Google Sheet for topics that are still pending. It sends each topic to our AI writing team. The team writes the content, checks quality, and then a real person approves it by email. After approval, the sheet says posted and that row is finished.

Think of it like a magazine office: robots draft the article, a human editor clicks Approve - and every draft should still try to help, not just impress.

3. How you generate content (two easy ways)

You do not need to know coding. Pick the way that fits your day.

Way 1 - Use the web interface (right now)

1. Open the app link you were given (the web page).
2. Type what you want to generate (the topic or brief) in the box.
3. Choose brand / content type if asked.
4. Click the Generate button.

That is it.

The full AI pipeline (research -> plan -> write -> review) usually takes about 1-2 minutes. When it finishes, the draft appears on the same page so you can read, copy, or download it.

Tip: Leave the page open while it runs. Do not close the tab until you see the result.

Way 2 - Use the Google Sheet (daily at 9 AM)

You will get a Google Sheet link. Use it like this:

1. Open the sheet.
2. In the topic column, write the topic name (what you want written).
3. In the status column, put Pending.

Example:

topic | status
How to start digital marketing for a startup | Pending

What happens next (automatic):

- Every day at 9 AM, the system looks for rows with status Pending.
- It runs the full writing pipeline for those topics.
- When done, it updates the generated_topic (or content) column with the result.
- You get an alert for approval (usually by email) so a human can say yes or no.

You do not need to click Generate for sheet topics - filling topic + Pending is enough. The 9 AM run does the rest.

4. Proper steps (start to finish)

Read this list in order. This is how the whole system works.

Step A - Clock rings (example: 9:00 AM)

n8n starts automatically (or someone runs it by hand).

Step B - Read Google Sheet

n8n opens the sheet and finds rows where:

- Status = `pending`

Example row:

Topic | Brand | Status
How to start digital marketing for a startup | Futuristix | pending

Step C - Call the AI writing API

n8n sends that topic to the LangGraph API (our backend).

This starts the multi AI agent pipeline.

Step D - Shared notebook (important idea)

All agents share one notebook called state (state.py for engineers).

- Each agent reads what it needs from the notebook

- Each agent writes its answer back into the same notebook
- The next agent uses that answer

Like one shared folder everyone updates in turn.

Step E - The five agents (in this exact order)

LangGraph is the traffic controller. It decides who goes next (graphing + routing).

```
1. Manager   (manager.py)
2. Research  (research.py)
3. Strategy  (strategy.py)
4. Writer    (writer.py)
5. Review    (review.py)
     |
     +-- if quality is low -> back to Writer (max 3 times)
```

1) Manager

- Checks the request
- Picks the brand (example: Futuristix)
- Understands if this should be article / email / social, etc.
- Writes brand details into the shared notebook

2) Research

- Searches the web / news / videos for useful facts
- Collects sources and number-style stats
- Writes the research pack into the notebook

3) Strategy (three special jobs)

Strategy plans "how to write," using three helpers:

A. SEO helper (`seo_service`)

- Creates keyword lists: primary, secondary, and long-tail
- Uses the research results and computer "meaning scores"
- Gives each keyword a score using things like:
  - how often it appears in research (frequency)
  - how close it is to the topic meaning (semantic similarity)
  - search intent (info vs buy vs browse)
  - match to customer pain points
  - brand / authority style fit
- Then it re-ranks keywords (best ones go first)

B. Hashtag helper

- Suggests hashtags (mainly useful for social)

C. Citation helper

- Cleans up source names / links for the article

Strategy also builds the outline (heading plan) and the writing blueprint.

4) Writer

- Reads the blueprint + research from the notebook
- Writes the full content
- Uses packaging helpers (formatter, json_builder) so the answer is neat for the sheet/API
- Hands the draft to Review

5) Review

- Scores the content out of 100
- Target score = 95
- If score is below 95 -> send notes back to Writer and try again
- This Writer ? Review loop can run up to 3 times

- Then the pipeline stops and returns the best result it has

Step F - Back to n8n + Google Sheet update

When the AI team finishes, n8n:

1. Puts the full article into the sheet column Content generated
2. Saves the score
3. Changes status from pending -> pending review

Step G - Human approval by email

n8n sends an email alert to a person:

- Topic
- Score
- Preview / link to content
- Buttons or reply: Approve / Reject

Step H - What happens next?

If human Approves

1. Sheet status becomes posted
2. This row's whole loop stops OK

If human Rejects

1. Content must be created again (n8n calls the AI API again)
2. Inside the AI team, Writer ? Review can still loop up to 3 times
3. If after recreates / retries it is still not good enough -> alert a developer ASAP so they can check and fix the system

5. Status meanings (sheet)

Status | Simple meaning
Pending (or pending) | Waiting. Not written yet. Put this after you add a **topic**.
pending review | Written. Waiting for human yes/no. Result is in **generated_topic** / content.
posted | Approved. Finished.
rejected | Human said no. Will recreate or need developer help.

Flow:

```
pending
   -> pending review   (content filled + email sent)
        -> posted                 (approve -> stop)
        -> recreate -> try again   (reject)
             -> still bad -> alert developer
```

6. Simple architecture picture

```
9 AM clock
    |
    v
[ n8n ] ---- reads ----> [ Google Sheet: pending topics ]
    |
    | calls API
    v
[ LangGraph traffic controller ]
    |
    +--> Manager
    +--> Research
    +--> Strategy (SEO + Hashtags + Citations)
    +--> Writer (+ formatter / json packer)
    +--> Review  --(if <95)--> Writer again (max 3)
    |
```

```
    v
[ n8n ] ---- writes content + "pending review" ----> [ Google Sheet ]
    |
    v
[ Email to human ]
    |
    +-- Approve --> status = posted --> STOP
    +-- Reject  --> recreate (up to 3) --> else ALERT DEVELOPER
```

7. Shared notebook example

Imagine the notebook after each agent:

1. Manager adds: brand = Futuristix, tone = practical
2. Research adds: 10 articles + 5 stats
3. Strategy adds: top keywords + outline + citations
4. Writer adds: full article text
5. Review adds: score = 96, status = pass

Then n8n copies the article text into the sheet.

8. Brands (who we write as)

Brand | Simple idea | Ending ask example
Futuristix | AI / growth for businesses | Book an AI Discovery Call
GTIB | Buying / selling companies | Schedule an Advisory Call
Kinvo | Childcare | Book a Consultation
MPM | Property | Contact Property Advisor
GCB | Engineering | Contact Our Engineering Team

9. What content types exist?

Type | Simple meaning
Article / Blog | Long guide or blog post
Email | Email campaign text
SEO content | Article plus keyword notes
Social | LinkedIn / short posts

10. Quality score in plain words

Review checks:

Check | Meaning
Content quality | Clear and useful?
SEO | Right search words?
Brand match | Sounds like our brand?
Structure | Good beginning / middle / end?
Facts | Claims backed by sources?
Call to action | Clear next step?

Goal: 95+. Retries inside AI: up to 3.

11. One full day example

8:55 AM

Sheet row: topic ready, status = pending

9:00 AM

n8n starts -> finds pending row -> calls AI API

9:08 AM (example timing)

Manager -> Research -> Strategy (SEO/hashtags/citations) -> Writer -> Review
(maybe Writer fixes twice if score was low)

9:10 AM

Sheet: Content generated filled, status = pending review
Email: "Please approve Futuristix article (score 95)"

10:00 AM

You click Approve -> status = posted -> stop

If you Reject

AI writes again (up to 3 tries). If still wrong -> developer gets an urgent alert.

12. Who does what?

Job | Who
Add topics to sheet (**topic** + **Pending**) | Marketing / ops
Use web page (type topic -> **Generate**) | Marketing / ops
Run at 9 AM, fill **generated_topic**, send approval alert | n8n
Research, plan, write, score | AI agents in this project
Approve / reject | Human
Fix system if rejects keep failing | Developer

13. Common questions

Q: Do I need to know code?

A: No. Use the web page (type topic -> Generate), or the sheet (topic + Pending) and the approval alert.

Q: How long does the web Generate button take?

A: Usually about 1-2 minutes for the full pipeline.

Q: Does "posted" mean it is live on the website?

A: In this design it means "approved and finished in the workflow." Website publish can be a later extra step.

Q: Why 9 AM?

A: Just an example schedule. Your team can choose any time in n8n.

Q: What if the score is 90 but a human likes it?

A: Human can still approve. The 95 goal guides the AI; humans have final say.

Q: How should we judge content when approving?

A: Ask: "Would this actually help someone stuck on this problem, or is it only flashy?" Prefer learning and clear next steps over views and hype.

14. One-page cheat sheet

Purpose first: help readers solve real problems - not chase views.

On the web page: write what you want -> click Generate -> wait about 1-2 minutes.

On the Google Sheet: put the topic in topic, put Pending in status -> daily 9 AM run -> generated_topic fills in -> alert for approval.

1. 9 AM - n8n starts
2. Take Pending topics from Google Sheet
3. Call LangGraph / multi-agent API
4. Manager -> Research -> Strategy -> Writer -> Review
5. Everyone uses shared state notebook
6. Strategy = SEO + hashtags + citations (SEO re-ranks keywords)
7. Writer packages with formatter / json builder
8. Review loop max 3, target score 95
9. Sheet gets generated_topic / content + status pending review
10. Email / alert to human for approval
11. Approve -> posted -> stop
12. Reject -> recreate (up to 3) -> else alert developer

Non-tech guide. Engineers: see Technical.md / Technical.pdf for file paths and scoring weights.